

CMPS242 Homework #2 – Chapter #1 Exercises

Zayd Hammoudeh

October 10, 2017

Suppose that we have three colored boxes, r (red), b (blue), and g (green). Box r contains 3 apples, 4 oranges, and 3 limes; box b contains 1 apple, 1 orange, and 0 limes; box g contains 3 apples, 3 oranges, and 4 limes.

- (a) If a box is chosen at random with probabilities, $p(r) = 0.2$, $p(b) = 0.2$, and $p(g) = 0.6$, and a piece of fruit is removed from the box (with equal probability of selecting any of the items in the box), then what is the probability of selecting an apple?

Recall the Law of Total Probability establishes that a set of disjoint events, B_1, \dots, B_n , that partition the sample space, S , (i.e., $\bigcup_{i=1}^n B_i = S$ and $\forall_{i,j} B_i \cap B_j = \emptyset$), then for any event A in S :

$$\Pr[A] = \sum_{i=1}^n \Pr[B_i] * \Pr[A|B_i]$$

Hence, define B as the set of all boxes (i.e., red, blue, and green). Then, the probability of selecting an apple can be found via:

$$\Pr[\text{apple}] = \sum_{b_i \in B} \Pr[\text{apple}|b_i] * \Pr[b_i].$$

This can be rewritten as:

$$\Pr[\text{apple}] = \Pr[\text{apple}|r] * \Pr[r] + \Pr[\text{apple}|b] * \Pr[b] + \Pr[\text{apple}|g] * \Pr[g]$$

$$\Pr[\text{apple}] = 0.3 * 0.2 + 0.5 * 0.2 + 0.3 * 0.6$$

$$\Pr[\text{apple}] = \boxed{0.34}$$

- (b) If we observe that the selected fruit is in fact an orange, what is the probability that it came from the green box?

The goal of this question is to find the posterior probability, $\Pr[g|\text{orange}]$. The simplest way to do that is to use the priors, $\Pr[\text{orange}]$ and $\Pr[g]$ with the likelihood, $\Pr[\text{orange}|g]$. Hence:

$$\Pr[g|\text{orange}] = \frac{\Pr[\text{orange}|g] * \Pr[g]}{\Pr[\text{orange}]} \quad (1)$$

Using again the Law of Total Probability, the prior probability of selecting an orange, $\Pr[o]$, is:

$$\Pr[\text{orange}] = \Pr[\text{orange}|r] * \Pr[r] + \Pr[\text{orange}|b] * \Pr[b] + \Pr[\text{orange}|g] * \Pr[g]$$

$$\Pr[\text{orange}] = 0.4 * 0.2 + 0.5 * 0.2 + 0.3 * 0.6$$

$$\Pr[\text{orange}] = 0.36$$

This can then be substituted into the Eq. (1).

$$\Pr[g|\text{orange}] = \frac{\Pr[\text{orange}|g] * \Pr[g]}{\Pr[\text{orange}]}$$

$$\Pr[g|\text{orange}] = \frac{0.3 * 0.6}{0.36}$$

$$\Pr[g|\text{orange}] = \boxed{0.5}$$

Show that the mode (i.e., the maximum) of the Gaussian distribution (1.46) is given by μ . Similarly, show that the mode of the multivariate Gaussian (1.52) is given by μ .

(a) Show the maximum of the Gaussian distribution is given by μ .

The maximizing value of a strictly positive function is equal to the maximizing value of the logarithm of the function. Below, “ln” is applied to the Gaussian formula.

$$\begin{aligned}\mathcal{N}(x|\mu, \sigma) &:= \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ \frac{-(x-\mu)^2}{2\sigma^2} \right\} \\ \ln(\mathcal{N}(x|\mu, \sigma)) &= \ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ \frac{-(x-\mu)^2}{2\sigma^2} \right\} \right) \\ &= \ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) + \ln \left(\exp \left\{ \frac{-(x-\mu)^2}{2\sigma^2} \right\} \right) \\ &= \ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) + \frac{-(x-\mu)^2}{2\sigma^2}\end{aligned}$$

The derivative can be applied and set to 0 as shown in Eq. (2).

$$0 = \frac{d}{dx} \ln(\mathcal{N}(x|\mu, \sigma)) = 0 - \frac{2(x-\mu)}{2\sigma^2} \quad (2)$$

Eq. (2) can be simplified yielding the final result as shown in Eq. (3).

$$x = \mu \quad \square \quad (3)$$

(b) Show the mode of the multivariate Gaussian distribution is given by μ .

The multivariate Gaussian distribution is shown in Eq. (4). \mathbf{x} is an n -dimensional input vector while μ is the n -dimensional mean vector. Σ is the $n \times n$ covariance matrix.

$$\mathcal{N}(\mathbf{x}|\mu, \Sigma) := \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu) \right\} \quad (4)$$

Similar to the technique used in part (a), the maximizing (mode) value of a strictly positive function is also the maximizing value of the function's natural log. Hence, Eq. (4) becomes:

$$\ln \mathcal{N}(\mathbf{x}|\mu, \Sigma) = \ln \left(\frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu) \right\} \right) \quad (5)$$

$$= \ln \left(\frac{1}{(2\pi)^{\frac{n}{2}}} \right) + \ln \left(\exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu) \right\} \right) \quad (6)$$

$$= \ln \left(\frac{1}{(2\pi)^{\frac{n}{2}}} \right) - \frac{1}{2} (\mathbf{x} - \mu)^T \Sigma^{-1} (\mathbf{x} - \mu). \quad (7)$$

The gradient of an n -dimensional function, $f(\mathbf{x})$, is an n -dimensional vector as defined in Eq. (8).

$$\nabla f(\mathbf{x}) = \left\langle \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n} \right\rangle \quad (8)$$

A multivariate function is maximized when its gradient equals the zero vector (i.e., $\mathbf{0}$). Hence, take the derivative of Eq. (7) and set it equal to the zero vector, which yields Eq. (10).

$$\mathbf{0} = \frac{\partial}{\partial \mathbf{x}} \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \Sigma) = \mathbf{0} + \frac{\partial}{\partial \mathbf{x}} (\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \quad (9)$$

$$\mathbf{0} = \frac{\partial}{\partial \mathbf{x}} (\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \quad (10)$$

If $\mathbf{y} = \mathbf{x} - \boldsymbol{\mu}$ and $A = \Sigma^{-1}$, then Eq. (10) becomes:

$$\mathbf{0} = \frac{\partial}{\partial \mathbf{y}} \mathbf{y}^T A \mathbf{y} \quad (11)$$

This matches the well known derivative identity that:

$$\frac{\partial}{\partial \mathbf{y}} \mathbf{y}^T A \mathbf{y} = (A^T + A) \mathbf{y}.$$

Hence, Eq. (11) simplifies to:

$$\mathbf{0} = (A^T + A) \mathbf{y}.$$

Substituting back in for \mathbf{y} completes the proof via:

$$(A^T + A)(\mathbf{x} - \boldsymbol{\mu}) = \mathbf{0}$$

$$\mathbf{x} - \boldsymbol{\mu} = \mathbf{0}$$

$$\mathbf{x} = \boldsymbol{\mu}. \quad \square$$

By setting the derivatives of the log likelihood function (1.54) with respect to μ and σ^2 equal to zero, verify the results (1.55) and (1.56).

The log likelihood function (Equation 1.54) is defined on page 27 as:

$$\ln p(\mathbf{x}|\mu, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi) \quad (12)$$

(a) Verify Equation (1.55) that:

$$\mu_{ML} = \frac{1}{N} \sum_{n=1}^N x_n$$

Take the partial derivative with respect to μ :

$$\begin{aligned} \frac{\partial}{\partial \mu} \left(-\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi) \right) &= \frac{-1}{2\sigma^2} \left(2 * -1 * \sum_{n=1}^N (x_n - \mu) \right) - 0 - 0 \\ &= \frac{\sum_{n=1}^N (x_n - \mu)}{\sigma^2} \\ &= \frac{\sum_{n=1}^N (x_n)}{\sigma^2} - \frac{N * \mu}{\sigma^2}. \end{aligned}$$

The right side of the equation can be set equal to 0. The denominator can be multiplied out resulting in:

$$\begin{aligned} 0 &= \sum_{n=1}^N (x_n) - N * \mu \\ N * \mu &= \sum_{n=1}^N x_n \\ \mu &= \frac{1}{N} \sum_{n=1}^N x_n \quad \square \end{aligned}$$

(b) Verify Equation (1.56) that:

$$\sigma_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{ML})^2$$

Taking the partial derivative with respect to σ and using the chain rule as necessary, we get:

$$\begin{aligned} \frac{\partial}{\partial \sigma} \left(-\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi) \right) &= -\frac{1}{2} \sum_{n=1}^N (x_n - \mu)^2 * (-2 * \sigma^{-3}) - \frac{N}{2} \left(\frac{1}{\sigma^2} (2\sigma) \right) - 0 \\ &= \frac{\sum_{n=1}^N (x_n - \mu)^2}{\sigma^3} - \frac{N}{\sigma} \end{aligned}$$

This equation is set equal to 0 and everything multiplied by σ resulting in:

$$\begin{aligned}
0 &= \frac{\sum_{n=1}^N (x_n - \mu)^2}{\sigma^2} - N \\
N &= \frac{\sum_{n=1}^N (x_n - \mu)^2}{\sigma^2} \\
\sigma^2 &= \frac{\sum_{n=1}^N (x_n - \mu)^2}{N}.
\end{aligned}$$

The optimal value for μ was found in part (a) which can then be substituted into the previous equation yielding:

$$\sigma^2 = \frac{\sum_{n=1}^N (x_n - \mu_{ML})^2}{N}. \quad \square$$

Chapter #1, Problem #29

Consider an M -state discrete random variable, x , and use Jensen's inequality in the form (1.115) to show that the entropy of its distribution $p(x)$ satisfies $\mathbf{H}[x] \leq \ln M$.

Equation (1.115) in the textbook (page 56) states:

$$f\left(\sum_{i=1}^M \lambda_i x_i\right) \leq \sum_{i=1}^M \lambda_i f(x_i).$$

A probability distribution, $p(x)$, satisfies the two Jensen Inequality criteria of λ namely: $\forall_i \lambda_i \geq 0$ and $\sum_{i=1}^M (\lambda_i) = 1$. This transforms Jensen's Inequality to:

$$f\left(\sum_{i=1}^M p(x_i) x_i\right) \leq \sum_{i=1}^M p(x_i) f(x_i). \quad (13)$$

Recall that a random variable is a function, $X : \Omega \rightarrow \mathbb{R}$, where Ω is the sample space (in this case $|\Omega| = M$) and \mathbb{R} is the set of real numbers. For this proof, define

$$x_i = \frac{1}{p(x_i)}.$$

This can be substituted into Eq. (13) yielding:

$$f\left(\sum_{i=1}^M p(x_i) \frac{1}{p(x_i)}\right) \leq \sum_{i=1}^M p(x_i) f\left(\frac{1}{p(x_i)}\right)$$

$$f\left(\sum_{i=1}^M 1\right) \leq \sum_{i=1}^M p(x_i) f\left(\frac{1}{p(x_i)}\right)$$

$$f(M) \leq \sum_{i=1}^M p(x_i) f\left(\frac{1}{p(x_i)}\right).$$

Define $f(x) = \ln\left(\frac{1}{x}\right)$ yielding:

$$\ln\left(\frac{1}{M}\right) \leq \sum_{i=1}^M p(x_i) \ln p(x_i).$$

Multiplying both sides by -1 (which inverts the inequality) and using the definition of entropy $H[x] = -\sum p(x) \ln p(x)$ completes the proof.

$$-\ln\left(\frac{1}{M}\right) \geq -\sum_{i=1}^M p(x_i) \ln \frac{1}{p(x_i)}$$

$$\ln M \geq \mathbf{H}[x]. \quad \square$$

Evaluate the Kullback-Leibler divergence (1.113) between two Gaussians, $p(x) = \mathcal{N}(x|\mu, \sigma^2)$ and $q(x) = \mathcal{N}(x|m, s^2)$.

For two continuous probability distributions, p and q , the Kullback-Leibler divergence is defined as:

$$\text{KL}(p||q) := \int p(x) \ln \frac{p(x)}{q(x)} dx.$$

Let's simplify the logarithm first. Substituting for $p(x)$ and $q(x)$ yields:

$$\begin{aligned} \ln \left(\frac{p(x)}{q(x)} \right) &= \ln \left(\frac{\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{(x-\mu)^2}{2\sigma^2} \right)}{\frac{1}{\sqrt{2\pi s^2}} \exp \left(-\frac{(x-m)^2}{2s^2} \right)} \right) \\ &= \ln \left(\frac{s}{\sigma} \exp \left(-\frac{(x-\mu)^2}{2\sigma^2} + \frac{(x-m)^2}{2s^2} \right) \right) \\ &= \ln \left(\frac{s}{\sigma} \right) + \ln \left(\exp \left(-\frac{(x-\mu)^2}{2\sigma^2} + \frac{(x-m)^2}{2s^2} \right) \right) \\ &= \ln \left(\frac{s}{\sigma} \right) - \frac{(x-\mu)^2}{2\sigma^2} + \frac{(x-m)^2}{2s^2}. \end{aligned}$$

Substituting this back into the original equation yields:

$$\text{KL}(p||q) = \int p(x) \left(\ln \left(\frac{s}{\sigma} \right) - \frac{(x-\mu)^2}{2\sigma^2} + \frac{(x-m)^2}{2s^2} \right) dx.$$

Expand the equation for simplification.

$$\text{KL}(p||q) = \int p(x) \ln \left(\frac{s}{\sigma} \right) - p(x) \frac{(x^2 - 2 \cdot x \cdot \mu + \mu^2)}{2\sigma^2} + p(x) \frac{(x^2 - 2 \cdot x \cdot m + m^2)}{2s^2} dx. \quad (14)$$

Given $p(x) = \mathcal{N}(x|\mu, \sigma^2)$, a few useful identities are in Eq. (15), (16), and (17).

$$\int c \cdot p(x) dx = 1 \quad (15)$$

$$\int x \cdot p(x) dx = \mu \quad (16)$$

$$\int x^2 \cdot p(x) dx = \mu^2 + \sigma^2 \quad (17)$$

Using these identities, Eq. (14) simplifies to:

$$\begin{aligned} \text{KL}(p||q) &= \ln \left(\frac{s}{\sigma} \right) - \frac{(\mu^2 + \sigma^2 - 2 \cdot \mu^2 + \mu^2)}{2\sigma^2} + \frac{(\mu^2 + \sigma^2 - 2 \cdot \mu \cdot m + m^2)}{2s^2} \\ &= \ln \left(\frac{s}{\sigma} \right) - \frac{1}{2} + \frac{\sigma^2 + (\mu - m)^2}{2s^2} \end{aligned}$$

If desired, the fraction can be pulled out resulting in the final form in Eq. (18).

$$\text{KL}(p||q) = \frac{1}{2} \left(\ln \left(\frac{s^2}{\sigma^2} \right) - 1 + \frac{\sigma^2 + (\mu - m)^2}{s^2} \right) \quad (18)$$

Consider two variables \mathbf{x} and \mathbf{y} having joint distribution $p(\mathbf{x}, \mathbf{y})$, show that the differential entropy of this pair of variables satisfies

$$\mathbf{H}[\mathbf{x}, \mathbf{y}] \leq \mathbf{H}[\mathbf{x}] + \mathbf{H}[\mathbf{y}]$$

with equality if and only if \mathbf{x} and \mathbf{y} are statistically independent.

The joint entropy of two variables can be written in terms of their marginal entropies and their mutual information, $\mathbf{I}[\mathbf{x}, \mathbf{y}]$. This is shown in Eq. (19).

$$\mathbf{H}[\mathbf{x}, \mathbf{y}] = \mathbf{H}[\mathbf{x}] + \mathbf{H}[\mathbf{y}] - \mathbf{I}[\mathbf{x}; \mathbf{y}]. \quad (19)$$

Cover & Thomas show that mutual information has the relationship:

$$\mathbf{I}[\mathbf{x}, \mathbf{y}] \geq 0 \quad (20)$$

with equality only in the case that \mathbf{x} and \mathbf{y} are independent (see (2.90) on page 28 of the second edition of Cover & Thomas). Given Eq. (20) and mutual information's non-negativity, it is clear that:

$$\mathbf{H}[\mathbf{x}] + \mathbf{H}[\mathbf{y}] - \mathbf{I}[\mathbf{x}; \mathbf{y}] \leq \mathbf{H}[\mathbf{x}] + \mathbf{H}[\mathbf{y}].$$

This in turn can be combined with Eq. (19) yielding:

$$\mathbf{H}[\mathbf{x}, \mathbf{y}] \leq \mathbf{H}[\mathbf{x}] + \mathbf{H}[\mathbf{y}]. \quad (21)$$

Equality holds when $\mathbf{I}[\mathbf{x}; \mathbf{y}] = 0$, which only occurs when \mathbf{X} and \mathbf{Y} are independent per Cover & Thomas. \square

By applying Jensen's inequality (1.115) with $f(x) = \ln x$, show that the arithmetic mean of a set of real numbers is never less than their geometrical mean.

For a set of numbers, $\{x_1, \dots, x_M\}$, the arithmetic mean is defined as:

$$\frac{\sum_{i=1}^M x_i}{M}.$$

In contrast, for the same set of numbers, the geometric mean is defined as:

$$\left(\prod_{i=1}^M x_i \right)^{\frac{1}{M}}.$$

Jensen's Inequality (1.115) is defined as:

$$f\left(\sum_{i=1}^M \lambda_i x_i\right) \leq \sum_{i=1}^M \lambda_i f(x_i).$$

If we substitute for $f(x) = \ln x$, the inequality becomes:

$$\ln\left(\sum_{i=1}^M \lambda_i x_i\right) \leq \sum_{i=1}^M \lambda_i \ln(x_i).$$

For Jensen's Inequality to hold, λ must satisfy two conditions namely: $\lambda_i \geq 0$ and $\sum_{i=1}^M (\lambda_i) = 1$. An obvious satisfying case is $\lambda_i = 1/M$ for all $i = 1 \dots M$. This then changes the equation to:

$$\ln\left(\frac{\sum_{i=1}^M x_i}{M}\right) \leq \frac{1}{M} \sum_{i=1}^M \ln(x_i).$$

Using the properties of logarithms, the right side is transformable to a product via:

$$\ln\left(\frac{\sum_{i=1}^M x_i}{M}\right) \leq \frac{1}{M} \ln\left(\prod_{i=1}^M x_i\right).$$

Using another property of logarithms, the multiplying scalar, $\frac{1}{M}$ can be brought inside the logarithm as:

$$\ln\left(\frac{\sum_{i=1}^M x_i}{M}\right) \leq \ln\left(\left(\prod_{i=1}^M x_i\right)^{\frac{1}{M}}\right).$$

Both sides are then raised to the power of e completing the proof.

$$\frac{\sum_{i=1}^M x_i}{M} \leq \left(\prod_{i=1}^M x_i\right)^{\frac{1}{M}} \quad \square$$

Extra Problem #1

Prove that the sum of two convex functions is also convex.

A function, $f(x)$, is convex over an interval (a, b) if for every $x_1, x_2 \in (a, b)$ and $0 \leq \lambda \leq 1$,

$$f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2). \quad (22)$$

Given a second convex function, $g(x)$, and an additional function, $h(x) = f(x) + g(x)$, add the definition of convexity for f to the definition of convexity for g as shown in Eq. (23).

$$f(\lambda x_1 + (1 - \lambda)x_2) + g(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2) + \lambda g(x_1) + (1 - \lambda)g(x_2). \quad (23)$$

This can be rearranged due to the associativity and commutativity of real valued addition.

$$(f(\lambda x_1 + (1 - \lambda)x_2) + g(\lambda x_1 + (1 - \lambda)x_2)) \leq \lambda(f(x_1) + g(x_1)) + (1 - \lambda)(f(x_2) + g(x_2)).$$

Substituting using the definition of $h(x)$ proves convexity via:

$$h(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda h(x_1) + (1 - \lambda)h(x_2). \quad \square$$